

Ashish Gossain

AI / Machine Learning Engineer | Data Science

Delhi, India • +91 8368073848 • ashishgossain174@gmail.com • [LinkedIn](#) • [GitHub](#)

Computer Science graduate specialising in Data Science and Artificial Intelligence, with four internships building AI systems for healthcare, document processing, information retrieval and data engineering. Experience spans RAG pipelines, knowledge graphs, LLM applications and the data infrastructure behind them.

EDUCATION

SRM University, Delhi-NCR, Delhi, India

B.Tech. Computer Science & Engineering — Data Science & Artificial Intelligence | May 2026

CGPA: 8.37/10.00 | Final four-semester average: 9.03/10.00 (90.3%)

Coursework: Data Science (Predictive Analysis), Data Warehousing & Data Mining, Machine Learning, Artificial Intelligence, Big Data Analytics, Cyber Security, Computer Networks, Operating Systems, DBMS, Analysis & Design of Algorithms, Engineering Mathematics I–III

Vivekanand Public School, Delhi, India

Class XII — Science (CBSE) | 2022 | 75.50% Class X (CBSE) | 2020 | 87.66%

EXPERIENCE

FutureVerse — AI Engineer Intern

New Delhi (Onsite) · Jun 2026 – Present

- Worked on FV Orchestration, the platform that runs FutureVerse’s AI pipelines as composable, resumable skills in a shared workspace. Designed content-hash version identification so long-running LLM workflows restart from the last valid state instead of re-running expensive model calls, and altered content cannot retain its identifier — making every run reproducible and auditable.
- Built an RDF knowledge graph powering an adaptive learning engine: an ontology registry constrains permissible node and edge types, PostgreSQL backs storage, and validity periods let the system query what a student knew at any past point, so a wrong answer resolves to the concept at fault rather than a whole chapter.

Celebal Technologies — Data Science Intern (CSI)

Noida (Remote) · Jun – Aug 2025

- Built CareCompanion, a RAG medical assistant, as capstone of a mentor-evaluated internship; engineered a LangChain retrieval pipeline with all-MiniLM-L6-v2 embeddings and a FAISS vector store over ~200 clinical documents, passing retrieved context to Mistral-7B-Instruct-v0.3 for grounded, source-cited answers; deployed via Streamlit. [GitHub](#)

ADQVEST Capital — Data Analyst Intern

Chennai (Remote) · Jan – Mar 2025

- Built a Python ETL pipeline scraping product data from multiple e-commerce sites into MySQL, replacing a manual process.
- Wrote pre-ingestion validation reconciling inconsistent field formats across sources, repairing recoverable records and rejecting the rest before they reached the database.

Nocturne GmbH — Data Scientist Intern

Berlin (Remote) · Jul – Sep 2024

- Researched AI-driven clinical decision-support for neurological disease detection from retinal imaging.
- Curated and standardised ~2,000 patient retinal images from five sources differing in resolution, dimensions, annotation format and labelling.

PROJECTS

- **Talk-to-Your-Database (Text-to-SQL)** — Streamlit app turning plain-English questions into PostgreSQL via Qwen2.5-7B-Instruct (Hugging Face Inference API). An `is_safe()` layer in front of execution permits only SELECT and blocks INSERT/UPDATE/DELETE/DROP, so a generated query can never mutate the database. Connects to Supabase Postgres through psycopg2, renders an editable query with a plain-language explanation, and exports to CSV/PDF. [GitHub](#)
- **Invoice Extractor (Multilingual)** — Streamlit app using Gemini 1.5 Flash’s multimodal API to read invoice images (JPG/PNG) directly and answer questions across multiple languages, with no OCR preprocessing step. Deployed on Streamlit Cloud, with a local pdf2image/Poppler variant for PDFs. [Live demo](#) · [GitHub](#)
- **ASL to Text & Speech — Minor Project** — Real-time American Sign Language recognition. MediaPipe extracts 21 hand landmarks per frame, which feed a CNN classifier for gesture prediction; recognised characters are assembled into words and voiced through a text-to-speech module, giving a continuous signing-to-speech loop. [GitHub](#)
- **JARVIS — B.Tech Major Project (12 credits)** — Voice assistant integrating speech recognition, intent handling and text-to-speech into one conversational loop for desktop task automation.

TECHNICAL SKILLS

Programming: Python, SQL, CYPHER; C, Java and R at coursework level

Machine Learning: machine learning, deep learning, NLP, computer vision, predictive modelling, EDA, model evaluation

Data & Retrieval: pandas, NumPy, scikit-learn, RAG, vector embeddings, FAISS, LangChain, Hugging Face, RDF & knowledge graphs, Mistral, Qwen, Gemini, ColPali

Data Engineering: ETL, web scraping, data validation, Kafka, Hadoop, REST APIs, microservices, content hashing

Databases: PostgreSQL, MySQL, MongoDB, Redis, Supabase, Neo4j, RDF triple stores

Tools: Git, Jupyter, VS Code, Streamlit, MediaPipe, OpenCV, Mistral OCR, PaddleOCR, PyMuPDF, Microsoft Office

CERTIFICATIONS, LANGUAGES & LEADERSHIP

Certifications: [Data Science Job Simulation \(Forage / BCG X\)](#) · [AI & Machine Learning, Grade A+ \(Atos Prayas Foundation / ICT Academy\)](#) · [Complete A.I., ML & Data Science Bootcamp \(Udemy\)](#) · [Deep Learning for Developers \(Infosys Springboard\)](#)

Languages: [English — IELTS 7.5 \(CEFR C1\)](#) · [German — A2](#) · Hindi — native

Leadership: Unwind Connect and Cure (NGO) — [PR Head & Core Team Member](#) (2024 – Present) · volunteer since 2022; lead outreach and communications for the volunteer team